

# Holistic Specifications for Robust Programs

Sophia Drossopoulou<sup>1</sup>[0000–1111–2222–3333], Julian Mackay<sup>2</sup>[0000–1111–2222–3333],  
James Noble<sup>2</sup>[0000–1111–2222–3333], and Susan Eisenbach<sup>1</sup>[0000–0001–9072–6689]

<sup>1</sup> Imperial College London {scd,susan}@imperial.ac.uk

<sup>2</sup> Victoria University of Wellington {julian.mackay,kjx}@ecs.vuw.ac.nz

**Abstract** Functional specifications describe what program components *can* do — the *sufficient* conditions to invoke a component’s operations. Specifications of sufficient conditions are not enough to reason about components that interact with external components of possibly unknown provenance, as programs evolve over time. In this *open world* setting, we must also consider the *necessary* conditions, *i.e.* what are the conditions without which an effect will not happen.

In this paper we propose a language *Chainmail* for writing *holistic* specifications, fit for open world scenarios, that focus on necessary conditions (as well as sufficient conditions). We give a formal semantics whose core has been mechanised in the Coq proof assistant.

## 1 Introduction

Software guards our secrets, our money, our intellectual property, our reputation [?]. We entrust personal and corporate information to software which works in an *open* world, where it interacts with third party software of unknown provenance, possibly buggy and potentially malicious.

This means we need our software to be *robust*: to behave correctly even if used by erroneous or malicious third parties. We expect that our bank will only make payments from our account if instructed by us, or by somebody we have authorised, that space on a web given to an advertiser will not be used to obtain access to our bank details [?], or that a concert hall will not book the same seat more than once.

While language mechanisms such as constants, invariants, object capabilities [?], and ownership [?] make it *possible* to write robust programs, they cannot *ensure* that programs are robust. Ensuring robustness is difficult because it means different things for different systems: perhaps that critical operations should only be invoked with the requisite authority; perhaps that sensitive personal information should not be leaked; or perhaps that resource belonging to one user should not be consumed by another.g To ensure robustness, we need ways to specify what robustness means for a particular program, and ways to demonstrate that the particular program adheres to its specific robustness requirements.

In this paper we propose *Chainmail*, a specification language to express holistic specifications. The design of *Chainmail* was guided by the study of a sequence of examples from the object-capability literature and the smart contracts world: the membrane [?], the DOM [?,?], the Mint/Purse [?], the Escrow [?], the DAO [?,?] and ERC20 [?]. As we worked through the examples, we found a small set of language constructs

that let us write holistic specifications across a range of different contexts. In particular, *Chainmail* extends traditional program specification languages [?,?] with features which talk about:

**Permission:** Which objects may have access to which other objects; this is central since access to an object usually also grants access to the functions it provides.

**Control:** Which objects called functions on other objects; this is useful in identifying the causes of certain effects - eg funds can only be reduced if the owner called a payment function.

**Time:** What holds some time in the past, the future, and what changes with time,

**Space:** Which parts of the heap are considered when establishing some property, or when performing program execution; a concept related to, but different from, memory footprints and separation logics,

**Viewpoint:** Which objects and which configurations are internal to our component, and which are external to it; a concept related to the open world setting.

While many individual features of *Chainmail* can be found in other work, their power and novelty for specifying open systems lies in their careful combination. The contributions of this paper are:

- the design of the holistic specification language *Chainmail*,
- the semantics of *Chainmail*, and a Coq mechanisation of its core,
- the application of *Chainmail* to a sequence of examples.

The rest of the paper is organised as follows: first Section 3 presents the *Chainmail* specification language. Section 4 introduces the formal model underlying *Chainmail*, and then section 5 defines the semantics of *Chainmail*'s assertions. Section ?? shows how key points of exemplar problems can be specified in *Chainmail*, section ?? discusses our design, 6 considers related work, and section 7 concludes. We relegate various details to appendices.

## 2 Motivating Example: The Bank

As a motivating example, we consider a simplified banking application, with objects representing `Accounts` or `Banks`. As in [?], `Accounts` belong to `Banks` and hold money (balances); with access to two `Accounts` of the same `Bank` one can transfer any amount of money from one to the other. We give a traditional specification in Figure 1.

The PRE-condition of `deposit` requires that the receiver and the first argument (`this` and `src`) are `Accounts` and belong to the same bank, that the second argument (`amt`) is a number, and that `src`'s balance is at least `amt`. The POST-condition mandates that `amt` has been transferred from `src` to the receiver. The function `makeNewAccount` returns a fresh `Account` with the same bank, and transfers `amt` from the receiver `Account` to the new `Account`. Finally, the function `newAccount` when run by a `Bank` creates a new `Account` with corresponding amount of money in it.<sup>3</sup>

<sup>3</sup> Note that our very limited bank specification doesn't even have the concept of an account owner.

```

method deposit(src, amt)
PRE:  this,src:Account  $\wedge$  this $\neq$ src  $\wedge$  this.myBank=src.myBank  $\wedge$ 
      amt: $\mathbb{N}$   $\wedge$  src.balance $\geq$ amt
POST: src.balance=src.balancepre-amt  $\wedge$  this.balance=this.
      balancepre+amt

method makeNewAccount(amt)
PRE:  this:Account  $\wedge$  amt: $\mathbb{N}$   $\wedge$  this.balance $\geq$ amt
POST: this.balance=this.balancepre-amt  $\wedge$  fresh result  $\wedge$ 
      result: Account  $\wedge$  this.myBank=result.myBank  $\wedge$  result.
      balance=amt

method newAccount(amt)
PRE:  this:Bank
POST: result: Account  $\wedge$  result.myBank=this  $\wedge$  result.balance
      =amt

```

**Figure 1.** Functional specification of Bank and Account

With such a specification the code below satisfies its assertion. Assume that `acm_acc` and `auth_acc` are Accounts for the ACM and for a conference paper author respectively. The ACM's `acm_acc` has a balance of 10,000 before an author is registered, while afterwards it has a balance of 11,000. Meanwhile the `auth_acc`'s balance will be 500 from a starting balance of 1,500 (barely enough to buy a round of drinks at the conference hotel bar).

```

assume acm_acc,auth_acc: Account  $\wedge$  acm_acc.balance=10000  $\wedge$ 
      auth_acc.balance=1500
acm_acc.deposit(auth_acc,1000)
assert acm_acc.balance=11000  $\wedge$  auth_acc.balance=500

```

This reasoning is fine in a closed world, where we only have to consider complete programs, where all the code in our programs (or any other systems with which they interact) is under our control. In an open world, however, things are more complex: our systems will be made up of a range of components, many of which we do not control; and furthermore will have to interact with external systems which we certainly do not control. Returning to our author, say some time after registering by executing the `deposit` code above, they attempt to pay for a round at the bar. Under what circumstances can they be sure they have enough funds in their account?

To see the problem, what if the bank provided a `steal` method that emptied out every account in the bank into a thief's account. If this method existed and if it were somehow called between registering at the conference and going to the bar, then the author would find an empty account.

The critical problem is that a bank implementation including a `steal` method would meet the functional specification of the bank from fig. 1, so long as the methods `deposit`, `makeNewAccount`, and `newAccount` meet their specification.

One obvious solution would be to return to a closed-world interpretation of specifications: we interpret specifications such as fig. 1 as *exact* in the sense that only implementations that meet the functional specification exactly, *with no extra methods or behaviour*, are considered as suitable implementations of the functional specification. The problem is that this solution is far too strong: it would for example rule out a bank that during software maintenance was given a new method `count` that simply counted the number of deposits that had taken place, or a method `notify` to enable the bank to occasionally send notifications to its customers.

What we need is some way to permit bank implementations that send notifications to customers, but to forbid implementations of `steal`. The key here is to capture the (implicit) assumptions underlying fig. 1, and to provide additional specifications that capture those assumptions. The following three informal requirements prevent methods like `steal`:

1. An account's balance can be changed only if a client calls the `deposit` method with the account as the receiver or as an argument.
2. An account's balance can be changed only if a client has access to that particular account.
3. The Bank/Account component does not leak access to existing accounts or banks.

Compared with the functional specification we have seen so far, these requirements capture *necessary* rather than *sufficient* conditions: Calling the `deposit` method to gain access to an account is necessary for any change to that account taking place. The function `steal` is inconsistent with requirement (1), as it reduces the balance of an `Account` without calling the function `deposit`. However, requirement (1) is not enough to protect our money. We need to (2) to avoid an `Account`'s balance getting modified without access to the particular `Account`, and (3) to ensure that such accesses are not leaked.

We can express these requirements through *Chainmail* assertions. Rather than specifying the behaviour of particular methods when they are called, we write assertions that range across the entire behaviour of the Bank/Account module.

- $$\begin{aligned}
 (1) &\triangleq \forall a. [ a : \text{Account} \wedge \text{changes}\langle a.\text{balance} \rangle \longrightarrow \\
 &\quad \exists o. [ \langle o \text{ calls } a.\text{deposit}(\_, \_) \rangle \vee \langle o \text{ calls } \_.\text{deposit}(a, \_) \rangle ] ] \\
 (2) &\triangleq \forall a. \forall S : \text{Set}. [ a : \text{Account} \wedge \langle \text{will}\langle \text{changes}\langle a.\text{balance} \rangle \rangle \text{ in } S \rangle \longrightarrow \\
 &\quad \exists o. [ o \in S \wedge \text{external}\langle o \rangle \wedge \langle o \text{ access } a \rangle ] ] \\
 (3) &\triangleq \forall a. \forall S : \text{Set}. [ a : \text{Account} \wedge \langle \text{will}\langle \exists o. [ \text{external}\langle o \rangle \wedge \langle o \text{ access } a \rangle ] \rangle \text{ in } S \rangle \\
 &\quad \longrightarrow \exists o'. [ o' \in S \wedge \text{external}\langle o' \rangle \wedge \langle o' \text{ access } a \rangle ] ]
 \end{aligned}$$

In the above and throughout the paper, we use an underscore (`_`) to indicate an existentially bound variable whose value is of no interest.

Assertion (1) says that if an account's balance changes (`changes⟨a.balance⟩`), then there must be some client object `o` that called the `deposit` method with `a` as a receiver or as an argument (`⟨o calls _.deposit(_)⟩`).

Assertion (2) similarly constrains any possible change to an account's balance: If at some future point the balance changes ( $\text{will} \langle \text{changes} \langle \dots \rangle \rangle$ ), and if this future change is observed with the state restricted to the objects from  $S$  (*i.e.*  $\langle \dots \text{ in } S \rangle$ ), then at least one of these objects ( $o \in S$ ) is external to the Bank/Account system ( $\text{external} \langle o \rangle$ ) and has (direct) access to that account object ( $\langle o \text{ access } a \rangle$ ). Notice that while the change in the balance happens some time in the future, the external object  $o$  has access to  $a$  in the *current* state. Notice also, that the object which makes the call to `deposit` described in (1), and the object which has access to  $a$  in the current state described in (2) need not be the same: It may well be that the latter passes a reference to  $a$  to the former (indirectly), which then makes the call to `deposit`.

It remains to think about how access to an Account may be obtained. This is the remit of assertion (3): It says that if at some time in the future of the state restricted to  $S$ , some object  $o$  which is external has access to some account  $a$ , and if  $a$  exists in the current state, then in the current state some object from  $S$  has access to  $a$ . Where  $o$  and  $o'$  may, but need not, be the same object. And where  $o'$  has to exist and have access to  $a$  in the *current* state, but  $o$  need not exist in the current state – it may be allocated later.

A holistic specification for the bank account, then, would be our original sufficient functional specification from fig. 1 plus the necessary specifications (1)-(3) from above. This holistic specification permits an implementation of the bank that also provides `count` and `notify` methods, even though the specification does not mention either method. Critically, though, the *Chainmail* specification does not permit an implementation that includes a `steal` method. First, the `steal` method clearly changes the balance of every account in the bank, but assertion (1) requires that any method that changes the balance of any account must be called `deposit`. Second, the `steal` method changes the balance of every account in the system, and will do so without the caller having a reference to most of those accounts, thus breaching assertion (2).

Assertion (3) gives essential protection when dealing with foreign, untrusted code. When an Account is given out to untrusted third parties, assertion (3) guarantees that this Account cannot be used to obtain access to further Accounts. The ACM does not trust its authors, and certainly does not want to give them access to `acm_acc`, which contains all of the ACM's money. Instead, in order to receive money, it will pass a secondary account used for incoming funds, `acm_incoming`, into which an author will pay the fee, and from which the ACM will transfer the money back into the main `acm_acc`. Assertion (3) is crucial, because it guarantees that even malicious authors could not use knowledge of `acm_incoming` to obtain access to the main `acm_acc` or any other account.

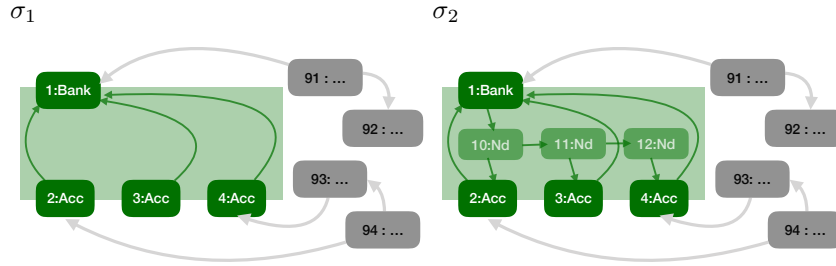
In summary, our necessary specifications are assertions that describe the behaviour of a module as observed by its clients. These assertions can talk about space, time and control, and thus go beyond class invariants, which can only talk about relations between the values of the fields of the objects. As our invariants are (usually) independent of concrete implementation details, they do not constrain the code to a specific implementation.

### 3 Chainmail Overview

In this Section we give a brief and informal overview of some of the most salient features of *Chainmail*—a full exposition appears in Section 5.

*Example Configurations* We will illustrate these features using the Bank/Account example from the previous Section. We use the runtime configurations  $\sigma_1$  and  $\sigma_2$  shown in the left and right diagrams in Figure 2. In both diagrams the rounded boxes depict objects: green for those from the Bank/Account component, and grey for the “external”, “client” objects. The transparent green rectangle shows which objects are contained by the Bank/Account component. The object at 1 is a Bank, those at 2, 3 and 4 are Accounts, and those at 91, 92, 93 and 94 are “client” objects which belong to classes different from those from the Bank/Account module.

Each configuration represents one alternative implementation of the Bank object. Configuration  $\sigma_1$  may arise from execution using a module  $M_{BA1}$ , where Account objects have a field `myBank` pointing to their Bank, and an integer field `balance`—the code can be found in appendix ?? Fig. ?? . Configuration  $\sigma_2$  may arise from execution using a module  $M_{BA2}$ , where Accounts have a `myBank` field, Bank objects have a `ledger` implemented though a sequence of Nodes, each of which has a field pointing to an Account, a field `balance`, and a field `next`—the code can be found in appendix ?? Figs. ?? and ?? .



**Figure 2.** Two runtime configurations for the Bank/Account example.

For the rest, assume variable identifiers  $b_1$ , and  $a_2$ – $a_4$ , and  $u_{91}$ – $u_{94}$  denoting objects 1, 2–4, and 91–94 respectively for both  $\sigma_1$  and  $\sigma_2$ . That is, for  $i=1$  or  $i=2$ ,  $\sigma_i(b_1)=1$ ,  $\sigma_i(a_2)=2$ ,  $\sigma_i(a_3)=3$ ,  $\sigma_i(a_4)=4$ ,  $\sigma_i(u_{91})=91$ ,  $\sigma_i(u_{92})=92$ ,  $\sigma_i(u_{93})=93$ , and  $\sigma_i(u_{94})=94$ .

*Classical Assertions* talk about the contents of the local variables (*i.e.* the topmost stack frame), and the fields of the various objects (*i.e.* the heap). For example, the assertion  $a_2.\text{myBank}=a_3.\text{myBank}$ , says that  $a_2$  and  $a_3$  have the same bank. In fact, this assertion is satisfied in both  $\sigma_1$  and  $\sigma_2$ , written formally as

$$\dots, \sigma_1 \models a_2.\text{myBank} = a_3.\text{myBank} \quad \dots, \sigma_2 \models a_2.\text{myBank} = a_3.\text{myBank}.$$

The term  $x:\text{ClassId}$  says that  $x$  is an object of class `ClassId`. For example

$$\dots, \sigma_1 \models a_2.\text{myBank} : \text{Bank}.$$

We support ghost fields  $[?, ?]$ , e.g.  $a_1.balance$  is a physical field in  $\sigma_1$  and a ghost field in  $\sigma_2$  since in MBA2 an Account does not store its balance (as can be seen in appendix ?? Fig. ??).

We also support the usual logical connectives, and so, we can express assertions such as

$$\forall a.[a : \text{Account} \longrightarrow a.myBank : \text{Bank} \wedge a.balance \geq 0].$$

*Permission: Access* Our first holistic assertion,  $\langle x \text{ access } y \rangle$ , asserts that object  $x$  has a direct reference to another object  $y$ : either one of  $x$ 's fields contains a reference to  $y$ , or the receiver of the currently executing method is  $x$ , and  $y$  is one of the arguments or a local variable. For example:

$$..., \sigma_1 \models \langle a_2 \text{ access } b_1 \rangle$$

Assuming that  $\sigma_1$  is executing the method body corresponding to the call  $a_2.deposit(a_3, 360)$ , we have

$$..., \sigma_1 \models \langle a_2 \text{ access } a_3 \rangle,$$

Namely, during execution of `deposit`, the object at  $a_2$  has access to the object at  $a_3$ , and could, if the method body chose to, call a method on  $a_3$ , or store a reference to  $a_3$  in its own fields.

Access is not symmetric, nor transitive:

$$\begin{aligned} ..., \sigma_1 &\not\models \langle a_3 \text{ access } a_2 \rangle, \\ ..., \sigma_2 &\models \langle a_2 \text{ access}^* a_3 \rangle, \quad ..., \sigma_2 \not\models \langle a_2 \text{ access } a_3 \rangle. \end{aligned}$$

*Control: Calls* The assertion  $\langle x \text{ calls } m.y(zs) \rangle$  holds in program states where a method on object  $x$  makes a method call  $y.m(zs)$  — that is it calls method  $m$  with object  $y$  as the receiver, and with arguments  $zs$ . For example,

$$..., \sigma_3 \models \langle x \text{ calls } a_2.deposit(a_3, 360) \rangle.$$

means that the receiver in  $\sigma_3$  is  $x$ , and the next statement to be executed is  $a_2.deposit(a_3, 360)$ .

*Space: In* The space assertion  $\langle A \text{ in } S \rangle$  establishes validity of  $A$  in a configuration restricted to the objects from the set  $S$ . For example, if object 94 is included in  $S_1$  but not in  $S_2$ , then we have

$$..., \sigma_1 \models \langle (\exists o. \langle o \text{ access } a_4 \rangle) \text{ in } S_1 \rangle \quad ..., \sigma_1 \not\models \langle (\exists o. \langle o \text{ access } a_4 \rangle) \text{ in } S_2 \rangle.$$

The set  $S$  in the assertion  $\langle A \text{ in } S \rangle$  is therefore *not* the footprint of  $A$ ; it is more like the *fuel*[?] given to establish that assertion. Note that  $..., \sigma \models \langle A \text{ in } S \rangle$  does not imply  $..., \sigma \models A$  nor does it imply  $..., \sigma \models \langle A \text{ in } S \cup S' \rangle$ . The other direction of the implication does not hold either.

*Time: Next, Will, Prev, Was* We support several operators from temporal logic:  $\langle \text{next} \langle A \rangle$ ,  $\text{will} \langle A \rangle$ ,  $\text{prev} \langle A \rangle$ , and  $\text{was} \langle A \rangle$  to talk about the future or the past in one or more number steps. The assertion  $\text{will} \langle A \rangle$  expresses that  $A$  will hold in one or more steps. For example, taking  $\sigma_4$  to be similar to  $\sigma_2$ , the next statement to be executed to be  $a_2.deposit(a_3, 360)$ , and  $M_{BA2} \circ ..., \sigma_4 \models a_2.balance = 60$ , and that  $M_{BA2} \circ ..., \sigma_4 \models a_4.balance \geq 360$

then

$$M_{BA2} \circ ..., \sigma_4 \models \text{will} \langle a_2.balance = 420 \rangle.$$

The *internal* module,  $M_{BA2}$  is needed for looking up the method body of `deposit`.

*Viewpoint: – External* The assertion  $\text{external}\langle x \rangle$  expresses that the object at  $x$  does not belong to the module under consideration. For example,

$$\begin{aligned} M_{AB2} \circ \dots, \sigma_2 &\models \text{external}\langle u_{92} \rangle, & M_{AB2} \circ \dots, \sigma_2 &\not\models \text{external}\langle a_2 \rangle, \\ M_{AB2} \circ \dots, \sigma_2 &\not\models \text{external}\langle b_1.\text{ledger} \rangle \end{aligned}$$

The *internal* module,  $M_{BA2}$ , is needed to judge which objects are internal or external.

*Change and Authority:* We have used  $\text{changes}\langle \rangle$  in our *Chainmail* assertions in section 2, as in  $\text{changes}\langle a.\text{balance} \rangle$ . Assertions that talk about change, or give conditions for change to happen are fundamental for security; the ability to cause change is called *authority* in [?]. We could encode change using the other features of *Chainmail*, namely, for any expression  $e$ :

$$\text{changes}\langle e \rangle \equiv \exists v. [e = v \wedge \text{next}\langle \neg(e = v) \rangle].$$

and similarly for assertions.

*Putting these together* We now look at some composite assertions which use several features from above. The assertion below says that if the statement to be executed is  $a_2.\text{deposit}(a_3, 60)$ , then the balance of  $a_2$  will eventually change:

$$M_{BA2} \circ \dots, \sigma_2 \models \langle .. \text{calls } a_2.\text{deposit}(a_3, 60) \rangle \longrightarrow \text{will}\langle \text{changes}\langle a_2.\text{balance} \rangle \rangle.$$

Now look deeper into space assertions,  $\langle A \text{ in } S \rangle$ : They allow us to characterise the set of objects which have authority over certain effects (here  $A$ ). In particular, the assertion  $\langle \text{will}\langle A \rangle \text{ in } S \rangle$  requires two things: i) that  $A$  will hold in the future, and ii) that the objects which cause the effect which will make  $A$  valid, are included in  $S$ . Knowing who has, and who has not, authority over properties or data is a fundamental concern of robustness [?]. Notice that the authority is a set, rather than a single object: quite often it takes *several objects in concert* to achieve an effect.

Now, consider assertions (2) and (3) from the previous section. They both have the form

$$\text{will}\langle \langle A \text{ in } S \rangle \rangle \longrightarrow P(S),$$

where  $P$  is some property over a set. These assertions say, that if ever in the future  $A$  becomes valid, and if the objects involved in making  $A$  valid are included in  $S$ , then  $S$  must satisfy  $P$ . Such assertions can be used to restrict whether  $A$  will become valid. Namely, if we have some execution which only involves objects which do not satisfy  $P$ , then we know that the execution will not ever make  $A$  valid.

*In summary*, in addition to classical logical connectors and classical assertions over the contents of the heap and the stack, our holistic assertions draw from some concepts from object capabilities ( $\langle \_ \text{access} \_ \rangle$  for permission;  $\langle \_ \text{calls} \_ \_ \rangle$  and  $\text{changes}\langle \_ \rangle$  for authority) as well as temporal logic ( $\text{will}\langle A \rangle$ ,  $\text{was}\langle A \rangle$  and friends), and the relation of our spatial connective ( $\langle A \text{ in } S \rangle$ ) with ownership and effect systems [?, ?, ?].

The next two sections discuss the semantics of *Chainmail*. Section 4 contains an overview of the formal model and section 5 focuses on the most important part of *Chainmail*: assertions.



## 4 Overview of the Formal foundations

We now give an overview of the formal model for *Chainmail*. In section 4.1 we introduce the shape of the judgments used to give semantics to *Chainmail*, while in section 4.2 we describe the most salient aspects of an underlying programming language used in *Chainmail*.

### 4.1 *Chainmail* judgments

Having outlined the ingredients of our holistic specification language, the next question to ask is: When does a module  $M$  satisfy a holistic assertion  $A$ ? More formally: when does  $M \models A$  hold?

Our answer has to reflect the fact that we are dealing with an *open world*, where  $M$ , our module, may be linked with *arbitrary untrusted code*. To model the open world, we consider pairs of modules,  $M \circledast M'$ , where  $M$  is the module whose code is supposed to satisfy the assertion, and  $M'$  is another module which exercises the functionality of  $M$ . We call our module  $M$  the *internal* module, and  $M'$  the *external* module, which represents potential attackers or adversaries.

We can now answer the question:  $M \models A$  holds if for all further, *potentially adversarial*, modules  $M'$  and in all runtime configurations  $\sigma$  which may be observed as arising from the execution of the code of  $M$  combined with that of  $M'$ , the assertion  $A$  is satisfied. More formally, we define:

$$M \models A \quad \text{if} \quad \forall M'. \forall \sigma \in \text{Arising}(M \circledast M'). [M \circledast M', \sigma \models A].$$

Module  $M'$  represents all possible clients of  $M$ . As it is arbitrarily chosen, it reflects the open world nature of our specifications.

The judgement  $M \circledast M', \sigma \models A$  means that assertion  $A$  is satisfied by  $M \circledast M'$  and  $\sigma$ . As in traditional specification languages [?,?], satisfaction is judged in the context of a runtime configuration  $\sigma$ ; but in addition, it is judged in the context of the internal and external modules. These are used to find abstract functions defining ghost fields as well as method bodies needed when judging validity of temporal assertions such as  $\text{will}(\_)$ .

We distinguish between internal and external modules. This offers two advantages: First, *Chainmail* includes the “external( $\circ$ )” assertion to require that an object belongs to the external module, as in the Bank Account’s assertion (2) and (3) in section 2. Second, we adopt a version of visible states semantics [?,?,?], treating all executions within a module as atomic. We only record runtime configurations which are *external* to module  $M$ , *i.e.* those where the executing object (*i.e.* the current receiver) comes from module  $M'$ . Execution has the form

$$M \circledast M', \sigma \rightsquigarrow \sigma'$$

where we ignore all intermediate steps with receivers internal to  $M$ . In the next section we shall outline the underlying programming language, and define the judgment  $M \circledast M', \sigma \rightsquigarrow \sigma'$  and the set  $\text{Arising}(M \circledast M')$ .

### 4.2 An underlying programming language, $\mathcal{L}_{\circ\circ}$

The meaning of *Chainmail* assertions is parametric with an underlying object-oriented programming language, with modules as repositories of code, classes with fields, meth-

ods and ghostfields, objects described by classes, a way to link modules into larger ones, and a concept of program execution.<sup>4</sup>

We have developed  $\mathcal{L}_{oo}$ , a minimal such object-oriented language, which we outline in this section. We describe the novel aspects of  $\mathcal{L}_{oo}$ , and summarise the more conventional parts, relegating full, and mostly unsurprising, definitions to Appendix ??,

Modules are central to  $\mathcal{L}_{oo}$ , as they are to *Chainmail*. As modules are repositories of code, we adopt the common formalisation of modules as maps from class identifiers to class definitions, c.f. Appendix, Def. ??. We use the terms module and component in an analogous manner to class and object respectively. Class definitions consist of field, method and ghost field declarations, c.f. Appendix, Def. ??.  $\mathcal{L}_{oo}$  is untyped – this reflects the open world, where we link with external modules which come without any guarantees. Method bodies are sequences of statements, which can be field read or field assignments, object creation, method calls, and return statements. Fields are private in the sense of C++: they can only be read or written by methods of the current class. This is enforced by the operational semantics, c.f. Fig. ??. We discuss ghost fields in the next section.

Runtime configurations,  $\sigma$ , contain all the usual information about an execution snapshot: the heap, and a stack of frames. The code to be executed is kept as part of the runtime configuration: Each frame consists of a continuation, `contn`, describing the remaining code to be executed by the frame, and a map from variables to values. Values are either addresses or sets of addresses; the latter are needed to deal with assertions which quantify over sets of objects, as e.g. (1) and (2) from section 2. We define *one-module* execution through a judgment of the form  $M, \sigma \rightsquigarrow \sigma'$  in the Appendix, Fig. ??.

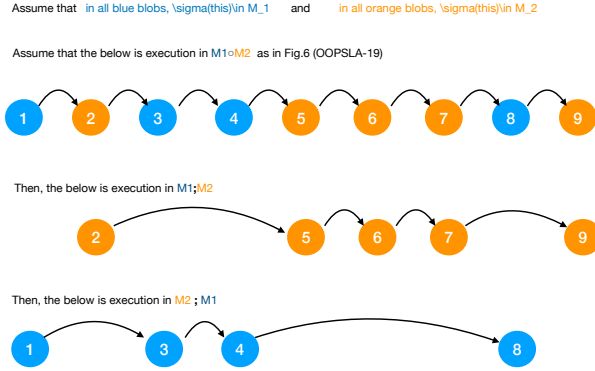
We define a module linking operator  $\circ$  so that  $M \circ M'$  is the union of the two modules, provided that their domains are disjoint, c.f. Appendix, Def. ??. As we said in section 4.1, we distinguish between the internal and external module, and treat execution of methods from the internal module as atomic. For this, we define *two-module execution* based on one-module execution as follows:

**Definition 1.** *Given runtime configurations  $\sigma, \sigma'$ , and a module-pair  $M \circ M'$  we define execution where  $M$  is the internal, and  $M'$  is the external module as below:*

- $M \circ M', \sigma \rightsquigarrow \sigma'$  if there exist  $n \geq 2$  and runtime configurations  $\sigma_1, \dots, \sigma_n$ , such that
  - $\sigma = \sigma_1$ , and  $\sigma_n = \sigma'$ .
  - $M \circ M', \sigma_i \rightsquigarrow \sigma'_{i+1}$ , for  $1 \leq i \leq n-1$
  - $\text{Class}(\text{this})_\sigma \notin \text{dom}(M)$ , and  $\text{Class}(\text{this})_{\sigma'} \notin \text{dom}(M)$ ,
  - $\text{Class}(\text{this})_{\sigma_i} \in \text{dom}(M)$ , for  $2 \leq i \leq n-2$

In the definition above,  $\text{Class}(x)_\sigma$  looks up the class of the object stores at  $x$ , c.f. Appendix, Def. ??. For example, for  $\sigma_4$  as in Section 3 whose next statement to be executed is `a2.deposit(a3, 360)`, we would have a sequence of configurations  $\sigma_{41}, \dots, \sigma_{4n}, \sigma_5$  so that the one-module execution gives  $M_{BA2}, \sigma_4 \rightsquigarrow \sigma_{41} \rightsquigarrow \sigma_{42} \dots \rightsquigarrow$

<sup>4</sup> We believe that *Chainmail* can be applied to any language with these features.



**Figure 3.** Illustrating Def. 1

$\sigma_{4n} \rightsquigarrow \sigma_5$ . This would correspond to an atomic evaluation in the two-module execution:  $M_{BA2} \circ M', \sigma_4 \rightsquigarrow \sigma_5$ .

We illustrate the definition in Fig. 3.

Two-module execution is related to visible states semantics [?] as they both filter configurations, with the difference that in visible states semantics execution is unfiltered and configurations are only filtered when it comes to the consideration of class invariants while two-module execution filters execution. The lemma below says that linking is associative and commutative, and preserves both one-module and two-module execution.

**Lemma 1 (Properties of linking).** *For any modules  $M, M', M''$ , and  $M'''$  and runtime configurations  $\sigma$ , and  $\sigma'$  we have:*

- $(M \circ M') \circ M'' = M \circ (M' \circ M'')$  and  $M \circ M' = M' \circ M$ .
- $M, \sigma \rightsquigarrow \sigma'$ , and  $M \circ M'$  is defined, implies  $M \circ M', \sigma \rightsquigarrow \sigma'$ .
- $M \circ M', \sigma \rightsquigarrow \sigma'$  implies  $(M \circ M'') \circ (M' \circ M'''), \sigma \rightsquigarrow \sigma'$ .

We can now answer the question as to which runtime configurations are pertinent when judging a module's adherence to an assertion. We define as *arising* configurations as those that can be reached by two-module execution, starting from any initial configuration. An initial configuration has a heap with only one object, of class `Object`, and only one frame, whose continuation may be arbitrary.

**Definition 2 (Initial and Arising Configurations).** *are defined as follows:*

- $\text{Initial}(\langle \psi, \chi \rangle)$ , if  $\psi$  consists of a single frame  $\phi$  with  $\text{dom}(\phi) = \{\text{this}\}$ , and there exists some address  $\alpha$ , such that  $[\text{this}]_\phi = \alpha$ , and  $\text{dom}(\chi) = \alpha$ , and  $\chi(\alpha) = (\text{Object}, \emptyset)$ .
- $\text{Arising}(M \circ M') = \{ \sigma \mid \exists \sigma_0. [\text{Initial}(\sigma_0) \wedge M \circ M', \sigma_0 \rightsquigarrow^* \sigma] \}$

## 5 Assertions

KJX: *Chainmail* Assertions consist of (pure) expressions  $e$ , classical assertions about the contents of heap and stack, the usual logical connectives, as well as our holistic concepts. In this section we focus on the novel holistic, features of *Chainmail* (permission, control, time, space, and viewpoint), as well as our wish to support some form of recursion while keeping the logic of assertions classical.

We now define the syntax and semantics of expressions and holistic assertions. The novel, holistic, features of *Chainmail* (permission, control, time, space, and viewpoint), as well as our wish to support some form of recursion while keeping the logic of assertions classical, introduced challenges, which we discuss in this section.

### 5.1 Satisfaction of Assertions - Access, Control, Space, Viewpoint

*Permission* expresses that an object has the potential to call methods on another object, and to do so directly, without help from any intermediary object. This is the case when the two objects are aliases, or the first object has a field pointing to the second object, or the first object is the receiver of the currently executing method and the second object is one of the arguments or a local variable.

**Definition 3 (Permission).** For any modules  $M, M'$ , variables  $x$  and  $y$ , we define

- $M \S M', \sigma \models \langle x \text{ access } y \rangle$  if  $\lfloor x \rfloor_\sigma$  and  $\lfloor y \rfloor_\sigma$  are defined, and
  - $\lfloor x \rfloor_\sigma = \lfloor y \rfloor_\sigma$ , or
  - $\lfloor x.f \rfloor_\sigma = \lfloor y \rfloor_\sigma$ , for some field  $f$ , or
  - $\lfloor x \rfloor_\sigma = \lfloor \text{this} \rfloor_\sigma$  and  $\lfloor y \rfloor_\sigma = \lfloor z \rfloor_\sigma$ , for some variable  $z$  and  $z$  appears in  $\sigma.\text{contn}$ .

In the last disjunct, where  $z$  is a parameter or local variable, we ask that  $z$  appears in the code being executed ( $\sigma.\text{contn}$ ). This requirement ensures that variables which were introduced into the variable map in order to give meaning to existentially quantified assertions, are not considered.

*Control* expresses which object is the process of making a function call on another object and with what arguments. The relevant information is stored in the continuation ( $\text{cont}$ ) on the top frame.

**Definition 4 (Control).** For any modules  $M, M'$ , variables  $x, y, z_1, \dots, z_n$ , we define:

- $M \S M', \sigma \models \langle x \text{ calls } y.m(z_1, \dots, z_n) \rangle$  if  $\lfloor x \rfloor_\sigma, \lfloor y \rfloor_\sigma, \lfloor z_1 \rfloor_\sigma, \dots, \lfloor z_n \rfloor_\sigma$  are defined, and
  - $\lfloor \text{this} \rfloor_\sigma = \lfloor x \rfloor_\sigma$ , and
  - $\sigma.\text{contn} = u.m(v_1, \dots, v_n); \_$  for some  $u, v_1, \dots, v_n$ , and
  - $\lfloor y \rfloor_\sigma = \lfloor u \rfloor_\sigma$ , and  $\lfloor z_i \rfloor_\sigma = \lfloor v_i \rfloor_\sigma$ , for all  $i$ .

Thus,  $\langle x \text{ calls } y.m(z_1, \dots, z_n) \rangle$  expresses the call  $y.m(z_1, \dots, z_n)$  will be executed next, and that the caller is  $x$ .

*Viewpoint* is about whether an object is viewed as belonging to the internal mode; this is determined by the class of the object.

**Definition 5 (Viewpoint).** For any modules  $M, M'$ , and variable  $x$ , we define


- $M \circ M', \sigma \models \text{external}\langle x \rangle$  if  $\lfloor x \rfloor_\sigma$  is defined and  $\text{Class}(\lfloor x \rfloor_\sigma)_\sigma \notin \text{dom}(M)$
- $M \circ M', \sigma \models \text{internal}\langle x \rangle$  if  $\lfloor x \rfloor_\sigma$  is defined and  $\text{Class}(\lfloor x \rfloor_\sigma)_\sigma \in \text{dom}(M)$

*Space* is about asserting that some property  $A$  holds in a configuration whose objects are restricted to those from a give set  $S$ . This way we can express that the objects from the set  $S$  have authority over the assertion  $A$ . In order to define validity of  $\langle A \text{ in } S \rangle$  in a configuration  $\sigma$ , we first define a restriction operation,  $\sigma \downarrow_S$  which restricts the objects from  $\sigma$  to only those from  $S$ .

**Definition 6 ( Restriction of Runtime Configurations).** The restriction operator  $\downarrow$  applied to a runtime configuration  $\sigma$  and a variable  $S$  is defined as follows:

- $\sigma \downarrow_S \triangleq (\psi, \chi'),$  if  $\sigma = (\psi, \chi)$ , and  $\text{dom}(\chi') = \lfloor S \rfloor_\sigma$ , and  $\forall \alpha \in \text{dom}(\chi'). \chi(\alpha) = \chi'(\alpha)$

For example, if we take  $\sigma_2$  from Fig. 2 in Section 2, and restrict it with some set  $S_4$  such that  $\lfloor S_4 \rfloor_{\sigma_2} = \{91, 1, 2, 3, 4, 11\}$ , then the restriction  $\sigma_2 \downarrow_{S_4}$  will look as on the right.



Note in the diagram above the dangling pointers at objects 1, 11, and 91 - reminiscent of the separation of heaps into disjoint subheaps, as provided by the  $*$  operator in separation logic [?]. The difference is, that in separation logic, the separation is provided through the assertions, where  $A * A'$  holds in any heap which can be split into disjoint  $\chi$  and  $\chi'$  where  $\chi$  satisfies  $A$  and  $\chi'$  satisfies  $A'$ . That is, in  $A * A'$  the split of the heap is determined by the assertions  $A$  and  $A'$  and there is an implicit requirement of disjointness, while in  $\sigma \downarrow_S$  the split is determined by  $S$ , and no disjointness is required.

We now define the semantics of  $\langle A \text{ in } S \rangle$ .

**Definition 7 ( Space ).** For any modules  $M, M'$ , assertions  $A$  and variable  $S$ , we define:

- $M \circ M', \sigma \models \langle A \text{ in } S \rangle$  if  $M \circ M', \sigma \downarrow_S \models A$ .

The set  $S$  in the assertion  $\langle A \text{ in } S \rangle$  is related to framing from implicit dynamic frames [?]: in an implicit dynamic frames assertion  $\text{acc } x.f * A$ , the frame  $x.f$  pre-prescribes which locations may be used to determine validity of  $A$ . The difference is that frames are sets of locations (pairs of address and field), while our  $S$ -es are sets of addresses. More importantly, implicit dynamic frames assertions whose frames are not large enough are badly formed, while in our work, such assertions are allowed and may hold or not, e.g.  $M_{BA2} \circ M', \sigma \models \neg \langle (\exists n. a_2.\text{balance} = n) \text{ in } S_4 \rangle$ .

## 5.2 Satisfaction of Assertions - Time

To deal with time, we are faced with four challenges: a) validity of assertions in the future or the past needs to be judged in the future configuration, but using the bindings from the current one, b) the current configuration needs to store the code being executed, so as to be able to calculate future configurations, c) when considering the future, we do not want to observe configurations which go beyond the frame currently at the top of the stack, d) there is no "undo" operator to deterministically enumerate all the previous configurations.

We discuss challenge a) in some more detail: take an assertion  $\text{will}\langle x.f = 3 \rangle$ ; it is satisfied in the *current* configuration,  $\sigma$ , if in some *future* configuration,  $\sigma'$ , the field  $f$  of the object that is pointed at by  $x$  in the *current* configuration ( $\sigma$ ) has the value 3, even in that future configuration  $x$  denotes a different object ( $[x]_\sigma \neq [x]_{\sigma'}$ ). To address this, we define an auxiliary concept:  $\triangleleft$  the adaptation of one runtime configuration to the scope of another.  $\sigma \triangleleft \sigma'$  adapts the second configuration to the top frame's view of the former: it returns a new configuration whose stack has the top frame as taken from  $\sigma$  and where the `contn` has been consistently renamed, while the heap is taken from  $\sigma'$ . This allows us to interpret expressions in  $\sigma'$  but with the variables bound according to  $\sigma$ ; *e.g.* we can obtain that value of  $x$  in configuration  $\sigma'$  even if  $x$  was out of scope in  $\sigma'$ .

**Definition 8 (Adaptation).** For runtime configurations  $\sigma, \sigma'$ :

- $\sigma \triangleleft \sigma' \triangleq (\phi'' \cdot \psi', \chi')$  if
  - $\sigma = (\phi \cdot \_, \_)$ ,  $\sigma' = (\phi' \cdot \psi', \chi')$ , and
  - $\phi = (\text{contn}, \beta)$ ,  $\phi' = (\text{contn}', \beta')$ ,  $\phi'' = (\text{contn}'[zs/zs'], \beta[zs' \mapsto \beta'(zs)])$ , where
  - $zs = \text{dom}(\beta)$ ,  $zs'$  is a set of variables with the same cardinality as  $zs$ , and
  - all variables in  $zs'$  are fresh in  $\beta$  and in  $\beta'$ .

That is, in the new frame  $\phi''$  from above, we keep the same continuation as from  $\sigma'$  but rename all variables with fresh names  $zs'$ , and combine the variable map  $\beta$  from  $\sigma$  with the variable map  $\beta'$  from  $\sigma'$  while avoiding names clashes through the renaming  $[zs' \mapsto \beta'(zs)]$ . The consistent renaming of the continuation allows the correct modelling of execution (as needed, for the semantics of nested time assertions, as *e.g.* in  $\text{will}\langle x.f = 3 \wedge \text{will}\langle x.f = 5 \rangle \rangle$ ).

Having addressed challenge a) we turn our attention to the remaining challenges: We address challenge b) by storing the remaining code to be executed in `contn` in each frame. We address challenge c) by only taking the top of the frame when considering future executions. Finally, we address challenge d) by considering only configurations which arise from initial configurations, and which lead to the current configuration.

**Definition 9 (Time Assertions).** For any modules  $M, M'$ , and assertion  $A$  we define

- $M \models M', \sigma \models \text{next}\langle A \rangle$  if  $\exists \sigma'. [ M \models M', \phi \rightsquigarrow \sigma' \wedge M \models M', \sigma \triangleleft \sigma' \models A ]$ ,  
and where  $\phi$  is so that  $\sigma = (\phi \cdot \_, \_)$ .
- $M \models M', \sigma \models \text{will}\langle A \rangle$  if  $\exists \sigma'. [ M \models M', \phi \rightsquigarrow^* \sigma' \wedge M \models M', \sigma \triangleleft \sigma' \models A ]$ ,  
and where  $\phi$  is so that  $\sigma = (\phi \cdot \_, \_)$ .

- $M \circ M', \sigma \models \text{prev}\langle A \rangle$  if  $\forall \sigma_1, \sigma_2. [ \text{Initial}\langle \sigma_1 \rangle \wedge M \circ M', \sigma_1 \rightsquigarrow^* \sigma_2 \wedge M \circ M', \sigma_2 \rightsquigarrow \sigma \longrightarrow M \circ M', \sigma \triangleleft \sigma_2 \models A ]$
- $M \circ M', \sigma \models \text{was}\langle A \rangle$  if  $\forall \sigma_1, \dots, \sigma_n. [ \text{Initial}\langle \sigma_1 \rangle \wedge \sigma_n = \sigma \wedge \forall i \in [1..n). M \circ M', \sigma_i \rightsquigarrow \sigma_{i+1} \longrightarrow \exists j \in [1..n-1). M \circ M', \sigma \triangleleft \sigma_j \models A ]$

In general,  $\text{will}\langle \langle A \text{ in } S \rangle \rangle$  is different from  $\langle \text{will}\langle A \rangle \text{ in } S \rangle$ . Namely, in the former assertion,  $S$  must contain the objects involved in reaching the future configuration as well as the objects needed to then establish validity of  $A$  in that future configuration. In the latter assertion,  $S$  need only contain the objects needed to establish  $A$  in that future configuration. For example, revisit Fig. 1, and take  $S_1$  to consist of objects 1, 2, 4, 93, and 94, and  $S_2$  to consist of objects 1, 2, 4. Assume that  $\sigma_5$  is like  $\sigma_1$ , that the next call in  $\sigma_5$  is a method on  $u_{94}$ , whose body obtains the address of  $a_4$  (by making a call on 93 to which it has access), and the address of  $a_2$  (to which it has access), and then makes the call  $a_2.\text{deposit}(a_4, 360)$ . Assume also and that  $a_4$  holds more than 360. Then

$$\begin{aligned} M_{BA1} \circ \dots, \sigma_5 &\models \langle \text{will}\langle \text{changes}\langle a_2.\text{balance} \rangle \rangle \text{ in } S_1 \rangle \\ M_{BA1} \circ \dots, \sigma_5 &\not\models \langle \text{will}\langle \text{changes}\langle a_2.\text{balance} \rangle \rangle \text{ in } S_2 \rangle \\ M_{BA1} \circ \dots, \sigma_5 &\models \text{will}\langle \langle \text{changes}\langle a_2.\text{balance} \rangle \text{ in } S_2 \rangle \rangle \end{aligned}$$

### 5.3 Properties of Assertions

We define equivalence of assertions in the usual way: assertions  $A$  and  $A'$  are equivalent if they are satisfied in the context of the same configurations and module pairs – *i.e.*

$$A \equiv A' \quad \text{if} \quad \forall \sigma. \forall M, M'. [ M \circ M', \sigma \models A \text{ if and only if } M \circ M', \sigma \models A' ].$$

We can then prove that the usual equivalences hold, *e.g.*  $A \vee A' \equiv A' \vee A$ , and  $\neg(\exists x.A) \equiv \forall x.(\neg A)$ . Our assertions are classical, *e.g.*  $A \wedge \neg A \equiv \text{false}$ , and  $M \circ M', \sigma \models A$  and  $M \circ M', \sigma \models A \rightarrow A'$  implies  $M \circ M', \sigma \models A'$ . This desirable property comes at the loss of some expected equivalences, *e.g.*, in general,  $e = \text{false}$  and  $\neg e$  are not equivalent. More in Appendix ??.

### 5.4 Modules satisfying assertions

Finally, we define satisfaction of assertions by modules: A module  $M$  satisfies an assertion  $A$  if for all modules  $M'$ , in all configurations arising from executions of  $M \circ M'$ , the assertion  $A$  holds.

**Definition 10.** For any module  $M$ , and assertion  $A$ , we define:

$$M \models A \quad \text{if} \quad \forall M'. \forall \sigma \in \text{Arising}(M \circ M'). M \circ M', \sigma \models A$$

## 6 Related Work

*Behavioural Specification Languages* Hatchliff et al. [?] provide an excellent survey of contemporary specification approaches. With a lineage back to Hoare logic [?], Meyer’s

Design by Contract [?] was the first popular attempt to bring verification techniques to object-oriented programs as a “whole cloth” language design in Eiffel. Several more recent specification languages are now making their way into practical and educational use, including JML [?], Spec# [?], Dafny [?] and Whiley [?]. Our approach builds upon these fundamentals, particularly Leino & Shulte’s formulation of two-state invariants [?], and Summers and Drossopoulou’s Considerate Reasoning [?]. In general, these approaches assume a closed system, where modules can be trusted to cooperate. In this paper we aim to illustrate the kinds of techniques required in an open system where modules’ invariants must be protected irrespective of the behaviour of the rest of the system.

*Defensive Consistency* In an open world, we cannot rely on the kindness of strangers: rather we have to ensure our code is correct regardless of whether it interacts with friends or foes. Attackers “*only have to be lucky once*” while secure systems “*have to be lucky always*” [?]. Miller [?,?] defines the necessary approach as **defensive consistency**: “*An object is defensively consistent when it can defend its own invariants and provide correct service to its well behaved clients, despite arbitrary or malicious misbehaviour by its other clients.*” Defensively consistent modules are particularly hard to design, to write, to understand, and to verify: but they have the great advantage that they make it much easier to make guarantees about systems composed of multiple components [?].

*Object Capabilities and Sandboxes.* *Capabilities* as a means to support the development of concurrent and distributed system were developed in the 60’s by Dennis and Van Horn [?], and were adapted to the programming languages setting in the 70’s [?]. *Object capabilities* were first introduced [?] in the early 2000s, and many recent studies manage to verify safety or correctness of object capability programs. Google’s Caja [?] applies sandboxes, proxies, and wrappers to limit components’ access to *ambient* authority. Sandboxing has been validated formally: Maffeis et al. [?] develop a model of JavaScript, demonstrate that it obeys two principles of object capability systems and show how untrusted applications can be prevented from interfering with the rest of the system. Recent programming languages and web systems [?,?,?] including Newspeak [?], Dart [?], Grace [?,?] and Wyvern [?] have adopted the object capability model.

*Verification of Object Capability Programs* Murray made the first attempt to formalise defensive consistency and correctness [?]. Murray’s model was rooted in counterfactual causation [?]: an object is defensively consistent when the addition of untrustworthy clients cannot cause well-behaved clients to be given incorrect service. Murray formalised defensive consistency very abstractly, over models of (concurrent) object-capability systems in the process algebra CSP [?], without a specification language for describing effects, such as what it means for an object to provide incorrect service. Both Miller and Murray’s definitions are intensional, describing what it means for an object to be defensively consistent.

Drossopoulou and Noble [?,?] have analysed Miller’s Mint and Purse example [?] and discussed the six capability policies as proposed in [?]. In [?], they sketched a specification language, used it to specify the six policies from [?], showed that several



possible interpretations were possible, and uncovered the need for another four further policies. They also sketched how a trust-sensitive example (the escrow exchange) could be verified in an open world [?]. Their work does not support the concepts of control, time, or space, as in *Chainmail*, but it offers a primitive expressing trust.

Sawsey et al. [?] have deployed powerful theoretical techniques to address similar problems: They show how step-indexing, Kripke worlds, and representing objects as state machines with public and private transitions can be used to reason about object capabilities. Devrise have demonstrated solutions to a range of exemplar problems, including the DOM wrapper (replicated in our section ??) and a mashup application. Their distinction between public and private transitions is similar to the distinction between internal and external objects.

More recently, Swasey et al. [?] designed OCPL, a logic for object capability patterns, that supports specifications and proofs for object-oriented systems in an open world. They draw on verification techniques for security and information flow: separating internal implementations (“high values” which must not be exposed to attacking code) from interface objects (“low values” which may be exposed). OCPL supports defensive consistency (they use the term “robust safety” from the security community [?]) via a proof system that ensures low values can never leak high values to external attackers. This means that low values *can* be exposed to external code, and the behaviour of the system is described by considering attacks only on low values. They use that logic to prove a number of object-capability patterns, including sealer/unsealer pairs, the caretaker, and a general membrane.

Schaefer et al. [?] have recently added support for information-flow security using refinement to ensure correctness (in this case confidentiality) by construction. By enforcing encapsulation, all these approaches share similarity with techniques such as ownership types [?,?], which also protect internal implementation objects from accesses that cross encapsulation boundaries. Banerjee and Naumann demonstrated that these systems enforce representation independence (a property close to “robust safety”) some time ago [?].

*Chainmail* differs from Swasey, Schaefer’s, and Devriese’s work in a number of ways: They are primarily concerned with mechanisms that ensure encapsulation (aka confinement) while we abstract away from any mechanism via the  $\text{external}\langle \rangle$  predicate. They use powerful mathematical techniques which the users need to understand in order to write their specifications, while the *Chainmail* users only need to understand first order logic and the holistic operators presented in this paper. Finally, none of these systems offer the kinds of holistic assertions addressing control flow, change, or temporal operations that are at the core of *Chainmail*’s approach.

KJX: need to talk about VerX

BLAH BLAH BLAH  
 BLAH BLAH BLAH  
 BLAH BLAH BLAH  
 BLAH BLAH BLAH  
 BLAH BLAH BLAH  
 BLAH BLAH BLAH  
 BLAH BLAH BLAH

BLAH BLAH BLAH  
BLAH BLAH BLAH  
BLAH BLAH BLAH  
BLAH BLAH BLAH  
BLAH BLAH BLAH  
BLAH BLAH BLAH  
BLAH BLAH BLAH  
BLAH BLAH BLAH

## 7 Conclusion

In this paper we have motivated the need for holistic specifications, presented the *Chainmail* specification language for writing such specifications, and shown how *Chainmail* can be used to give holistic specifications of key exemplar problems: the bank account, the wrapped DOM, the ERC20, and the DAO.

To focus on the key attributes of a holistic specification language, we have kept *Chainmail* simple, only requiring an understanding of first order logic. We believe that the holistic features (permission, control, time, space and viewpoint), are intuitive concepts when reasoning informally, and were pleased to have been able to provide their formal semantics in what we argue is a simple manner.